

# SliceDeque

A double-ended queue with contiguous elements

# About this talk

- My name is Gonzalo (@gnzlbq on Github)
- **Disclaimer:** I work @nvidia, these thoughts are my own, and do not necessarily reflect the views of my employer.
- I wrote a somewhat popular crate called “slice\_deque”
  - ~200k all-time downloads
  - Network I/O applications, e.g., Firefox’s QUIC (neqo)
  - Audio applications, e.g., minimp3
- The idea behind it is not new, but it is an idea worth sharing:  
<https://fgiesen.wordpress.com/2012/07/21/the-magic-ring-buffer/> (and others)

# std::VecDeque recap

```
let mut deq = VecDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
for _ in 0..3 { deq.pop_front(); }  
deq.extend(&[7, 8, 9]);
```

**O(1) push\_front/back**

**O(1) pop\_front/back**

# std::VecDeque recap

- `let mut deq = VecDeque::new();`  
`deq.extend(&[0, 1, 2, 3, 4, 5, 6]);` **O(1) push\_front/back**  
`for _ in 0..3 { deq.pop_front(); }` **O(1) pop\_front/back**  
`deq.extend(&[7, 8, 9]);`

## Stack

ptr
cap
head
tail

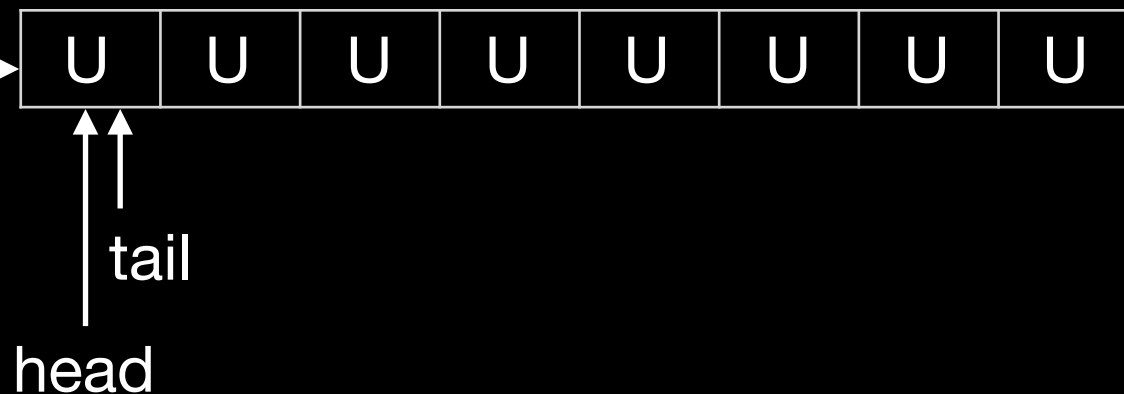
# std::VecDeque recap

- ```
let mut deq = VecDeque::new();
```
- `deq.extend(&[0, 1, 2, 3, 4, 5, 6]);` **O(1) push\_front/back**
  - `for _ in 0..3 { deq.pop_front(); }` **O(1) pop\_front/back**
  - `deq.extend(&[7, 8, 9]);`

**Stack**



**Heap**



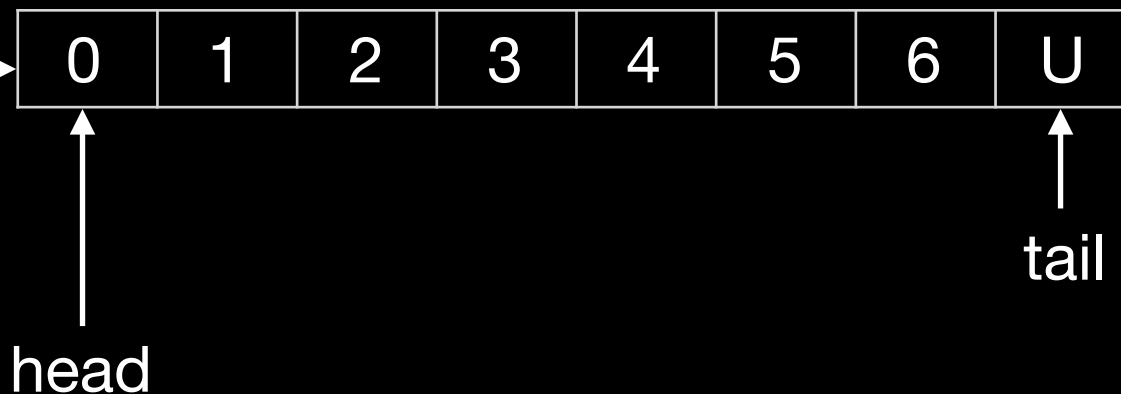
# std::VecDeque recap

- ```
let mut deq = VecDeque::new();
```
- `deq.extend(&[0, 1, 2, 3, 4, 5, 6]);` **O(1) push\_front/back**  
`for _ in 0..3 { deq.pop_front(); }` **O(1) pop\_front/back**  
`deq.extend(&[7, 8, 9]);`

**Stack**

ptr
cap
head
tail

**Heap**



# std::VecDeque recap

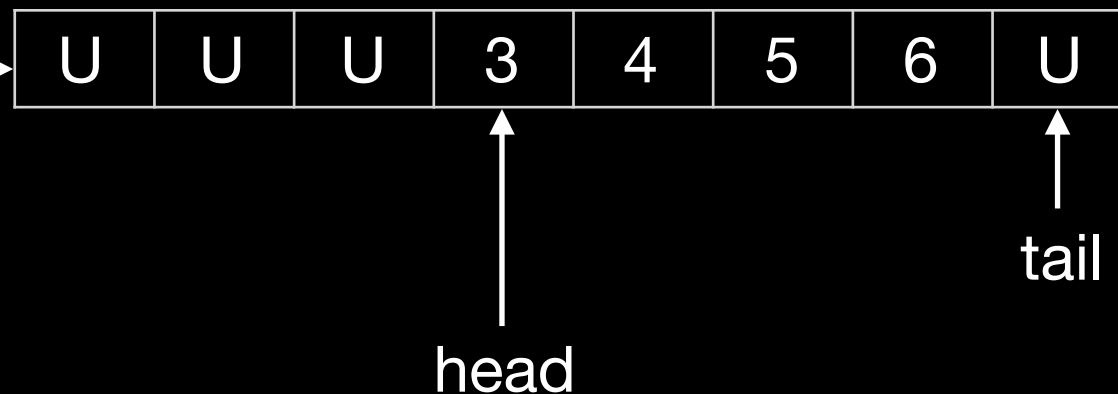
```
let mut deq = VecDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
• for _ in 0..3 { deq.pop_front(); }  
deq.extend(&[7, 8, 9]);
```

**O(1) push\_front/back**  
**O(1) pop\_front/back**

**Stack**

ptr
cap
head
tail

**Heap**



# std::VecDeque recap

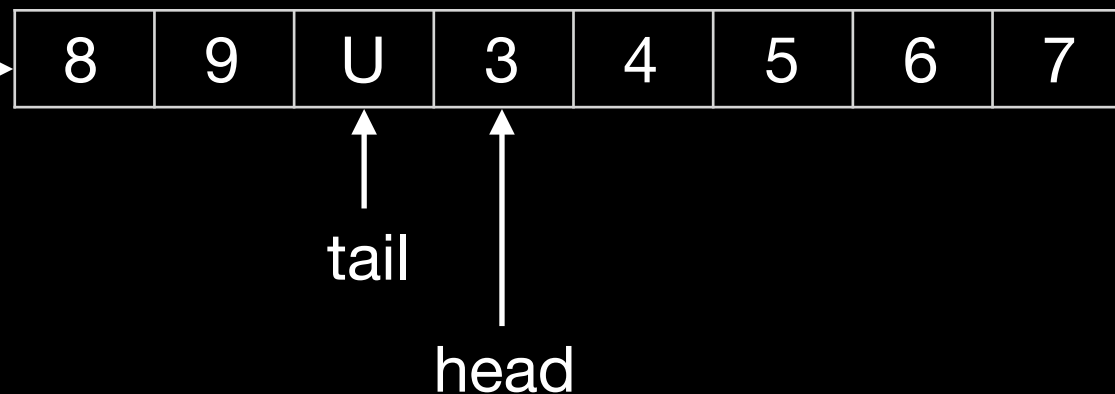
```
let mut deq = VecDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
for _ in 0..3 { deq.pop_front(); }  
● deq.extend(&[7, 8, 9]);
```

**O(1) push\_front/back**  
**O(1) pop\_front/back**

**Stack**

ptr
cap
head
tail

**Heap**

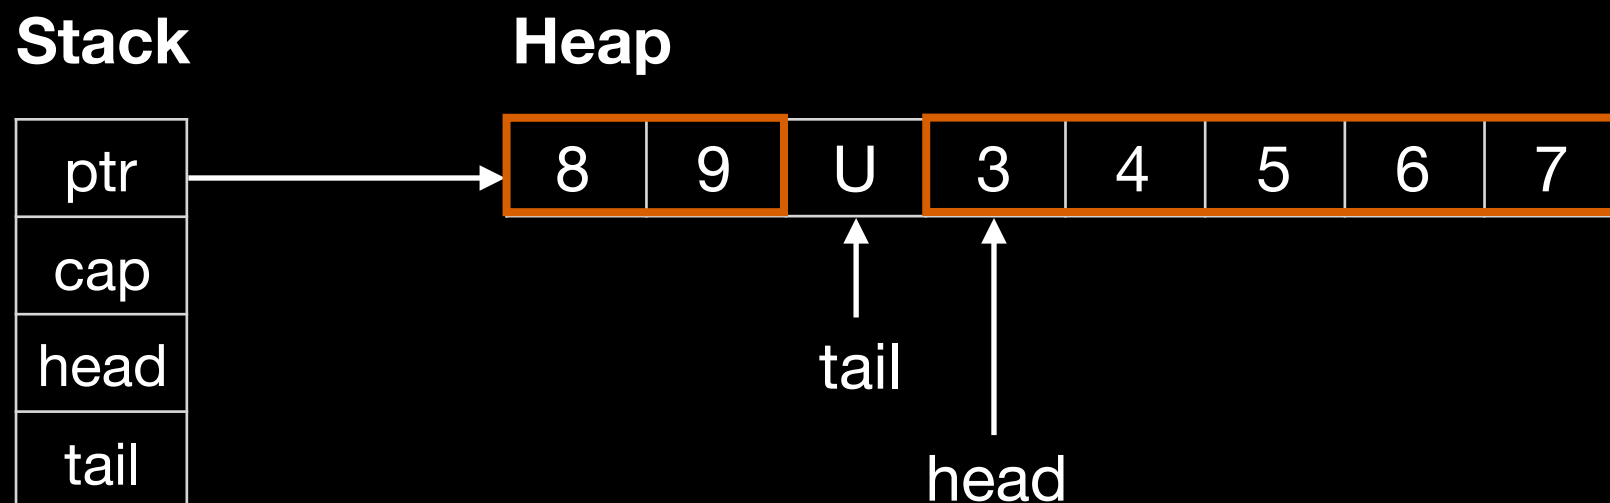




# std::VecDeque recap

```
let mut deq = VecDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
for _ in 0..3 { deq.pop_front(); }  
● deq.extend(&[7, 8, 9]);
```

**O(1) push\_front/back**  
**O(1) pop\_front/back**



- Empty if `head == tail`  $\Rightarrow$  `SZmax = cap - 1`
- **Efficiency**: accessing the  $i$ -th element is non-trivial (powers-of-two cap), hierarchical iterators
- **API**, `VecDeque::as_slices(&self) -> (&[T], &[T])`, requires some “duck tape” to interface with
  - This usually means that if you are processing bytes from the network via the Bytes trait, you need to batch the processing around the end of the deque
- Can we do better?

# Virtual memory 101

# Virtual memory 101

- OS gives each process its own private “virtual” address space, isolating it from all other processes (protection), and abstracting how “virtual” addresses are backed by “physical” memory.
- **Security:** CPU traps if process tries to access memory w/o appropriate rights
- **Efficiency:** OS can change how virtual memory is backed (RAM, NVM, swap)
- MMUs, MPUs, etc. provide hardware support for virtual memory & protection

# Virtual memory 101


- OS gives each process its own private “virtual” address space, isolating it from all other processes (protection), and abstracting how “virtual” addresses are backed by “physical” memory.
- **Security:** CPU traps if process tries to access memory w/o appropriate rights
- **Efficiency:** OS can change how virtual memory is backed (RAM, NVM, swap)
- MMUs, MPUs, etc. provide hardware support for virtual memory & protection

Physical memory



# Virtual memory 101

- OS gives each process its own private “virtual” address space, isolating it from all other processes (protection), and abstracting how “virtual” addresses are backed by “physical” memory.
- **Security:** CPU traps if process tries to access memory w/o appropriate rights
- **Efficiency:** OS can change how virtual memory is backed (RAM, NVM, swap)
- MMUs, MPUs, etc. provide hardware support for virtual memory & protection

Physical memory 

## Process A

allocate 4 linked list nodes

VM 

# Virtual memory 101

- OS gives each process its own private “virtual” address space, isolating it from all other processes (protection), and abstracting how “virtual” addresses are backed by “physical” memory.
- **Security:** CPU traps if process tries to access memory w/o appropriate rights
- **Efficiency:** OS can change how virtual memory is backed (RAM, NVM, swap)
- MMUs, MPUs, etc. provide hardware support for virtual memory & protection



## Process A

allocate 4 linked list nodes

delete one node



# Virtual memory 101

- OS gives each process its own private “virtual” address space, isolating it from all other processes (protection), and abstracting how “virtual” addresses are backed by “physical” memory.
- **Security:** CPU traps if process tries to access memory w/o appropriate rights
- **Efficiency:** OS can change how virtual memory is backed (RAM, NVM, swap)
- MMUs, MPUs, etc. provide hardware support for virtual memory & protection



## Process A

allocate 4 linked list nodes  
delete one node



## Process B

allocate a contiguous 3 node array



# SliceDeque

```
let mut deq = SliceDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
for _ in 0..3 { deq.pop_front(); }  
deq.extend(&[7, 8, 9]);
```



# SliceDeque

- `let mut deq = SliceDeque::new();`  
`deq.extend(&[0, 1, 2, 3, 4, 5, 6]);`  
`for _ in 0..3 { deq.pop_front(); }`  
`deq.extend(&[7, 8, 9]);`

## Stack

ptr
cap
head
tail

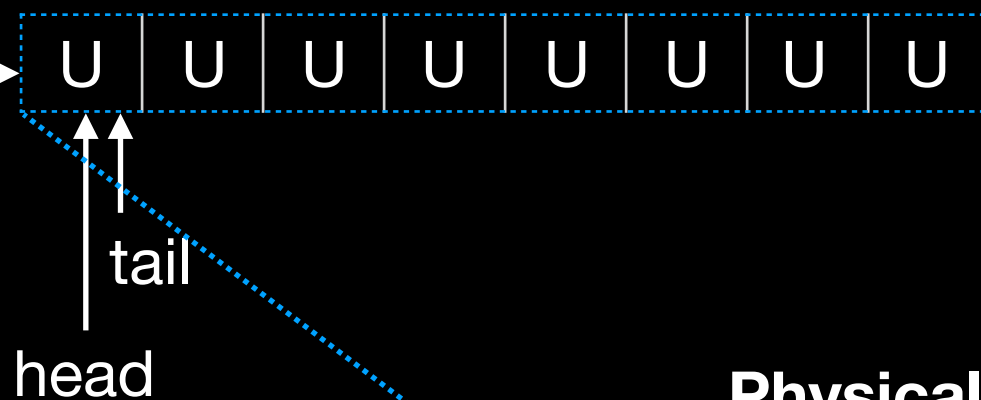
# SliceDeque

```
let mut deq = SliceDeque::new();  
● deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
  for _ in 0..3 { deq.pop_front(); }  
  deq.extend(&[7, 8, 9]);
```

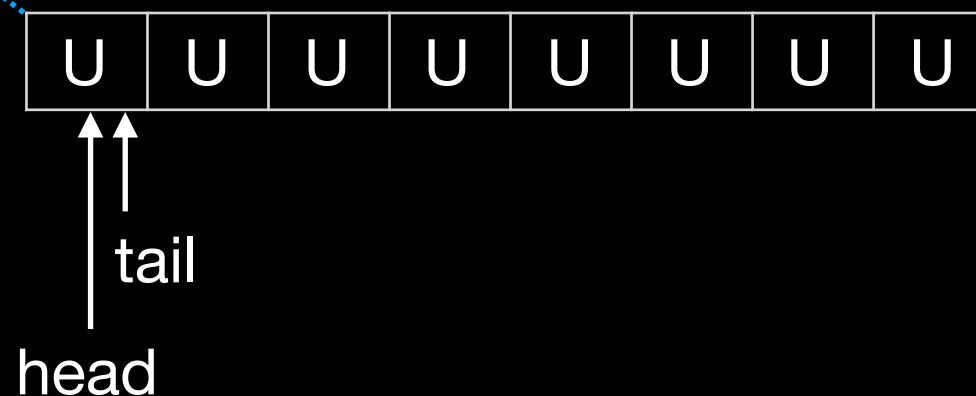
**Stack**

ptr
cap
head
tail

**Heap (virtual memory)**



**Physical memory**



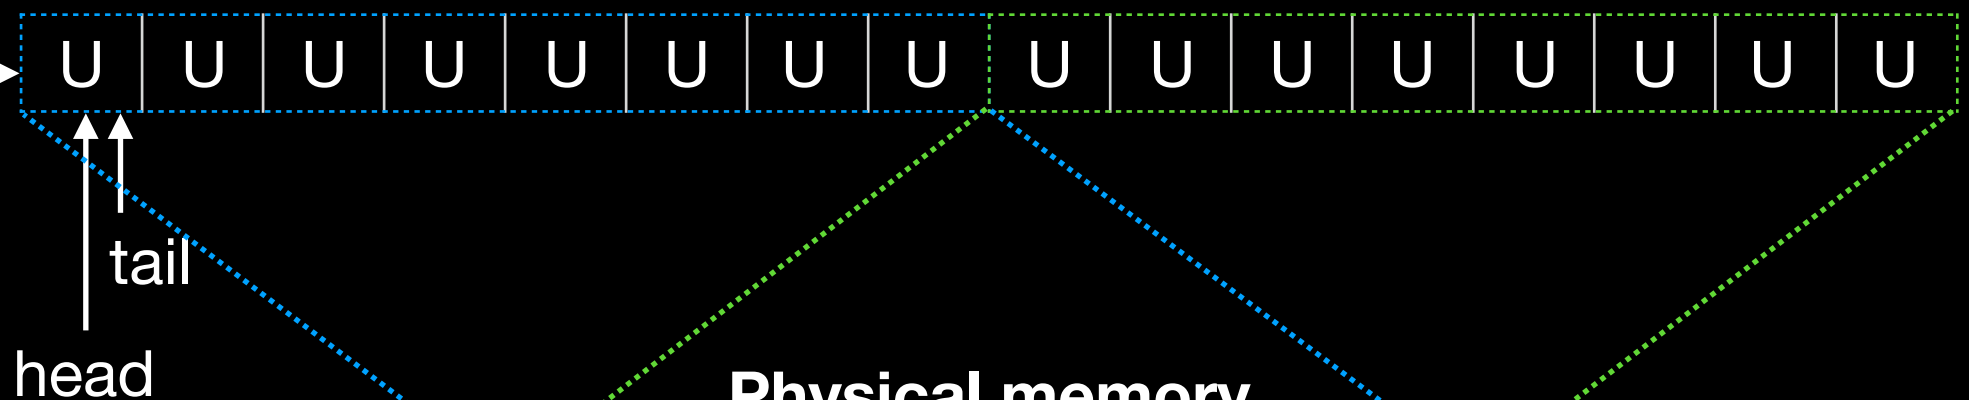
# SliceDeque

```
let mut deq = SliceDeque::new();  
● deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
  for _ in 0..3 { deq.pop_front(); }  
  deq.extend(&[7, 8, 9]);
```

**Stack**

ptr
cap
head
tail

**Heap (virtual memory)**



**Physical memory**



head  
tail

**Mirror the physical  
memory to an  
adjacent virtual  
memory region**

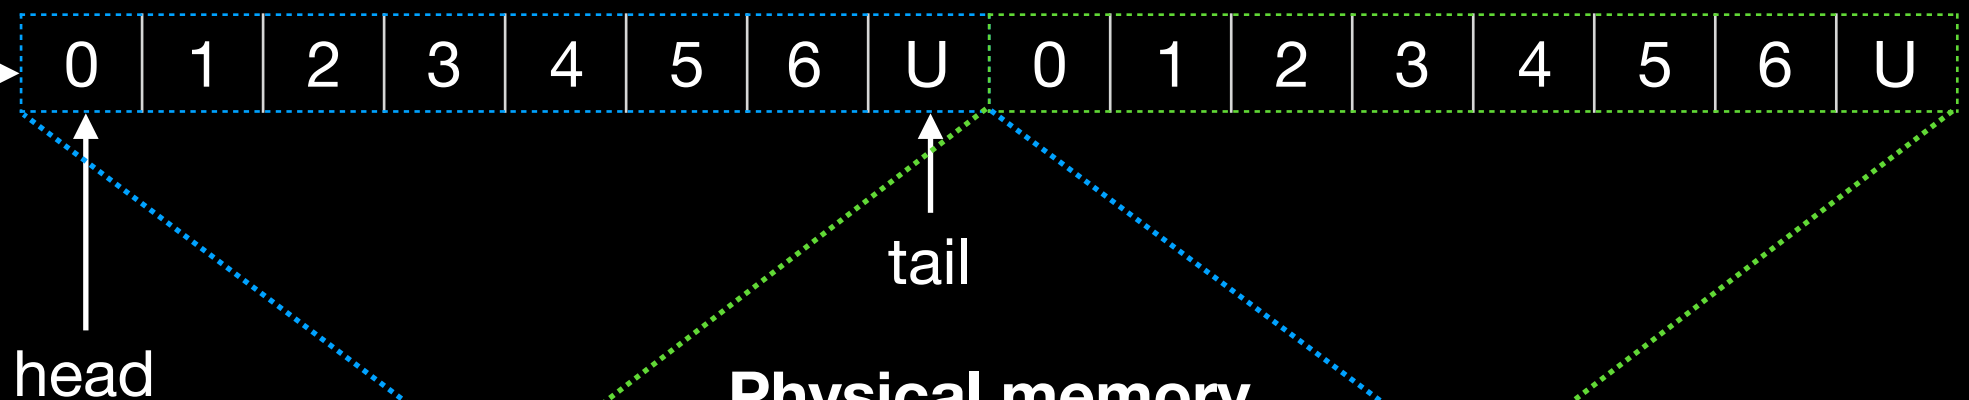
# SliceDeque

```
let mut deq = SliceDeque::new();  
● deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
  for _ in 0..3 { deq.pop_front(); }  
  deq.extend(&[7, 8, 9]);
```

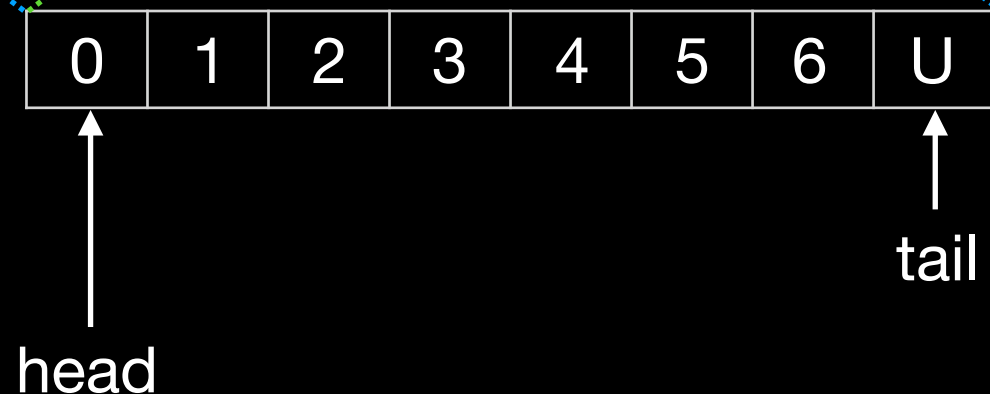
**Stack**

ptr
cap
head
tail

**Heap (virtual memory)**



**Physical memory**



# SliceDeque

```
let mut deq = SliceDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
• for _ in 0..3 { deq.pop_front(); }  
deq.extend(&[7, 8, 9]);
```

**Stack**

ptr
cap
head
tail

**Heap (virtual memory)**



head

tail

**Physical memory**



head

tail

# SliceDeque

```
let mut deq = SliceDeque::new();  
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);  
for _ in 0..3 { deq.pop_front(); }  
● deq.extend(&[7, 8, 9]);
```

**Stack**

ptr
cap
head
tail

**Heap (virtual memory)**



head

tail

**Physical memory**



tail

head

# SliceDeque

```
let mut deq = SliceDeque::new();
deq.extend(&[0, 1, 2, 3, 4, 5, 6]);
for _ in 0..3 { deq.pop_front(); }
• deq.extend(&[7, 8, 9]);
```

**In virtual memory, the deque elements are always contiguous!**

**Stack**

ptr
cap
head
tail

**Heap (virtual memory)**



**Physical memory**



- Simple indexing
- `&[T]` zero-copy

# Pros / Cons

- **Widely supported:** SliceDeque supports Linux, MacOSX, Windows, iOS, Android, FreeBSD, OpenBSD, Solaris, NetBSD, ... It is trivial to port it to any POSIX compliant OS.
- **Simple API:** interface directly with all APIs using `&[T]` (sorting, binary search, etc.).
- **Efficient:** simple indexing and iterators, vectorization, bulk processing, etc.
- **Requires platform support:** the “OS” (and hardware) must support virtual memory. Qol is platform dependent: great on MacOSX, ok-ish on Linux  $\geq 3.17$ , less ok-ish everywhere else: resources, race-conditions; improvements in [#10](#).
- **Bypasses the global memory allocator:** requiring multiple syscalls on growth and relocation, etc. There is a PR open that improves it.
- **Constrained minimum capacity:** the memory usage of SliceDeque is often constrained to multiples of the allocation granularity of the OS 4/8/64 kB, since the virtual to physical memory map works on memory pages.



# When does it make sense?

If the problem requires:

- a fixed-number of deques,
- the deques are long-lived,
- each deque is fixed-size,
- their size is large enough (at least 1 memory page),
- the target supports it,
- the resources required by the target are acceptable,

then using a SliceDeque ***might*** make sense.

**Examples:** I/O buffers for: audio-channels, executors, etc.

# That's all Folks!

- **Idea:** explicit virtual memory use enables “impossible” data-structures  
**others:** non-reallocating vectors, concurrent lock-free deques, etc.
- The details are at: [https://github.com/gnzlb主/slice\\_deque](https://github.com/gnzlb主/slice_deque)
- Great Rust-specific resources:
  - **Writing an OS in Rust** (x64) by Philipp Opperman: *Introduction to Paging* (<https://os.phil-opp.com/paging-introduction/>) and *Paging implementation* (<https://os.phil-opp.com/paging-implementation/>)
  - **RISC-V OS using Rust** by Stephen Marz: *Page-grained memory allocation* (<http://osblog.stephenmarz.com/ch3.html>) and *Memory-management unit* (<http://osblog.stephenmarz.com/ch3.2.html>)
- Thank you for your attention and I wish you a happy hacking!